

WASSERSTEIN AUTO-ENCODERS

MARZIEH MEHDIZADEH, JOSEPH D VIVIANO, JOHNATHAN GUYMONT

1. INTRODUCTION

Method to estimate the density distribution of data have been studied for a long time. Density estimation is indeed required in many applied science fields (e.g. physics, economics, insurance etc.). Formally, given an observed dataset $X \in \mathcal{X}$ the goal is to approximate the true distribution P_X with $\tilde{P}_X \in \mathcal{P}$ where \mathcal{P} is the set of all distribution of X . The model is then selected using some goodness of fit criteria, usually function of the likelihood, e.g. the selected model could be $\arg \max_{P \in \mathcal{P}} E \log P(X)$. A common method in practice is to make some hypothesis about the prior distribution (e.g. gamma, lognormal, ...) and find the parameters that gives the best fit to the data distribution for all of the selected prior distributions. The selected distribution is the one that gives the best fit over all the tested distribution. Note that model usually have an infinite number of possible parametrization θ and the number of possible model that are consistent with the definition of probability distribution is also infinite. Thus the approach is intractable and in practice a finite set of model to be tested is selected. This approach have huge drawback. If the best model is not in the set to be tested, then it is not possible to find the maximum. It is also likely to overfit and the specified model in these setting usually lack the capacity required to learn complex relation.

On the other hand, recently developed generative approaches, VAE and GAN, do not require hypothesis on the prior of the distribution of the data. These two approaches allow the use of highly flexible implicit model P_G to modeled complex underlying distribution from observed data. They are learn by comparing sample from the model P_G and observed data and they can be scaled up with SGD. Their setting comes with efficient regularizing method allowing good generalization. In these settings, the distribution is defined as $p(x) = \int p(x|z)p(z)dz$ where $p(z)$ is a fixed distribution and $p(x|z)$ is learned by comparing sample from the model $\tilde{x} \sim P_G$ and observed data x . This work present a summary of a newly developed approach to learn the (possibly random) mapping $p(x|z)$ presented in a recently published paper [Tolstikhin et al.(2017)] with a reproduction of the experiments. In particular, they use an optimal transport cost to derive a new family of auto-encoders - the Wasserstein auto-encoders.

2. PAPER SUMMARY

2.1. Motivation. (What is the problem)

In this paper, the authors introduce a new efficient algorithm called Wasserstein Auto-Encoders (WAE) to build a generative model of data distribution. This generative model is based on the optimal transport (OT) finding the minimum Wasserstein distance between the true data probability distribution (unknown distribution) P_X and the latent variable distribution P_G . In

other words, WAE minimizes a penalized form of the Wasserstein distance between P_X and P_G which leads to a regularizer that is different from the regularizer in VAE. This regularizer helps the posterior $P_G(X|Z)$ distribution of data points $X \in \mathcal{X}$ to match the prior P_Z of latent codes $Z \in \mathcal{Z}$.

2.2. Proposed approach. (Analysis)

First we list the main contributions of this paper as below:

- (1) The first and the main contribution is the family of Wasserstein Auto-Encoders (WAE) defined as above. The loss function of WAE algorithm is a combination of two terms: a reconstruction cost (e.g. the mean square error between the ground truth and the generated data) plus a regularizer that penalized according to the discrepancy measured by $D_Z(P_Z, Q_Z)$ between a prior P_Z on Z and a learned encoder Q_Z .
- (2) We use two MNIST and CelebA datasets to evaluate the WAE algorithm and we will observe that the generated samples are less blurry than VAE's generated samples.
- (3) Two different form of regularizer $D_Z(P_Z, Q_Z)$ will be introduced: One is based on the adversarial training in latent space \mathcal{Z} (Like GANs) so we denote it the **GAN-based** \mathcal{D}_Z and the other form is based on the maximum mean discrepancy denoted as **MMD-based** \mathcal{D}_Z .
- (4) Theoretical part containing a theorem that helps finding objective function for WAE.

How WAE works:

The WAE method tries to minimize the optimal transport cost $W_c(P_X, P_G)$ such the the decoder and encoder have the following tasks:

Decoder: It tries to reconstruct the encoded training examples with a good accuracy by using the cost function c , defined by $c(x, y) = ||x - y||_2^2$.

Encoder: It simultaneously does two tasks: It tries to match the encoded distribution of training examples Q_z to the prior P_Z measured by using one of the form of $\mathcal{D}_Z(Q_Z, P_Z)$ introduced above, and at the same time it makes sure that the latent codes $Z \in \mathcal{Z}$ are learning enough to reconstruct the encoded training examples. In ?? we give more explanation that how WAE is more efficient than VAE.

What is Wasserstein distance? Wasserstein Distance is a measure of the distance between two probability distributions. (It is also called Earth Mover's distance, short for EM distance, because informally it can be interpreted as moving piles of dirt that follow one probability distribution at a minimum cost to follow the other distribution.) The cost is quantified by the amount of dirt moved times the moving distance. We take the minimum one among the costs of all dirt moving solutions as the EM distance. In the definition of Wasserstein distance, the \inf (infimum, also known as greatest lower bound) indicates that we are only interested in the smallest cost.) More formally we define the Wasserstein distance as follows:

$$(1) \quad W_c(P_X, P_G) := \inf_{\Gamma \in \mathcal{P}(X \sim P_X, Y \sim P_G)} \mathbb{E}_{(X,Y) \sim \Gamma} [c(X, Y)],$$

where $c(x, y)$ is any measurable cost function and $\mathcal{P}(X \sim P_X, Y \sim P_G)$ is the set of all possible joint probability distribution between P_X and P_G .




FIGURE 1. In these figures circles true data points X for data space \mathcal{X} , triangles are the latent points Z from space \mathcal{Z} , squares represent the generated samples Y and arrows show the different conditional distributions. In both cases VAE (left) and WAE (right) we have to minimize the $c(x, y)$ and $\mathcal{D}_Z(P_Z, Q_Z)$ but there is a big difference that as we observe in left, VAE tries to encourage every single $Q(Z|X = x)$ shown by the red balls to match P_Z shown by the white ball, but the problem is that the red balls have intersection which leads the blurriness problem to the reconstructed results. On the other hand, WAE forces the expectation $Q_Z = \mathbb{E}_{P_X} [Q(Z|X)]$ with the green balls to match P_Z shown by white balls. If we define $\Gamma_{VAE}(Y|X) = P_G(Y|Z) + Q_{VAE}(Z|X)$ and $\Gamma_{WAE}(Y|X) = P_G(Y|Z) + Q_{WAE}(Z|X)$ which are shown by the red arrows, then we see that in the left picture the data point X are mapped to a same Y with a high probability while in the right picture data points X are mapped to Y a lower probability.

Why Wasserstein distance is better than JS or KL divergence used in VAE?

Even when two distributions are located in lower dimensional manifolds without overlaps, Wasserstein distance can still provide a meaningful and smooth representation of the distance in-between. Let's give an examples: If P and Q are two probability distributions and we define:

$$\begin{aligned} \forall (x, y) \in P, x = 0, y \sim U(0, 1) \\ \forall (x, y) \in Q, x = \theta, 0 \leq \theta \leq 1, y \sim U(0, 1) \end{aligned}$$

when $\theta \neq 0$.

$$D_{KL}(P||Q) = \sum_{x=0, y \sim U(0,1)} 1 \cdot \log(1/0) = +\infty,$$

$$D_{KL}(Q||P) = \sum_{x=\theta, y \sim U(0,1)} 1 \cdot \log(1/0) = +\infty,$$

$$D_{Js}(P, Q) = 1/2 \left(\sum_{X=0, y \sim U(0,1)} 1 \cdot \log \frac{1}{(1/2)} + \sum_{X=0, Y \sim U(0,1)} 1 \cdot \log \frac{1}{(1/2)} \right) = \log 2$$

and

$$W(P, Q) = |\theta|$$

But when $\theta = 0$, two distributions are fully overlapped:

$$D_{KL}(P||Q) = D_{KL}(Q||P) = D_{JS}(P, Q) = 0,$$

and

$$W(P, Q) = 0 = |\theta|$$

D_{KL} gives us infinity when two distributions are disjoint. The value of D_{JS} has sudden jump, not differentiable at $\theta = 0$. Only Wasserstein metric provides a smooth measure, which is super helpful for a stable learning process using gradient descents.

3. APPLICATION TO GENERATIVE MODELS

3.1. Tractable method for computing the Wasserstein distance $W_c(P_X, P_G)$. The goal of generative models is to approximate the distribution of the data P_X with a latent variable models P_G . The distribution is defined as $p_G(x) = \int p_G(x|z)p_Z(z)dz$ where $p_Z(z)$ is a fixed distribution (e.g. standard normal) and $p_G(x|z)$ is learned by comparing sample from the model $\tilde{x} \sim P_G$ and observed data x . To define P_G , a distribution P_Z is choose as a prior over the latent variable Z . Then Z is sampled from P_Z and \tilde{x} is obtain from the mapping $P_G : Z \rightarrow X$. The goal of the learning algorithm is to learn the mapping P_G such that equation (1) is minimized. The Wasserstein distance is intractable, so instead the algorithm also learn a conditional distribution $Q(Z|X)$ such that its expected value over X is equal to the prior P_Z , i.e.

$$p_G(x) = \int p_G(x|z)p_Z(z)dz = E_{X \sim P_X} \int p_G(x|z)q(z|x)dx$$

The reason why defining P_G like this is useful is because theorem 1 in [Tolstikhin et al.(2017)] state that minimizing (1) is equivalent to minimizing

$$(2) \quad W_c(P_X, P_G) = \inf_{Q: Q_Z = P_Z} E_{P_X} E_{Q(Z|X)}[c(X, G(Z))]$$

3.2. Regularizer. Minimizing the cost c between the data X and sample $G(Z)$ should force the model to learn a good mapping but Q_Z still need to be equal to P_Z for equality (2) to hold. For this reason, the loss function of WAE algorithm is a combination of a reconstruction cost (e.g. the mean square error between the ground truth and the generated data) plus a regularizer that penalized according to the discrepancy measured by $D_Z(P_Z, Q_Z)$ between a prior P_Z on Z and a learned encoder Q_Z . Adding this regularizer should ensure that Q_Z is close to P_Z . The main difference between VAE and WAE is this discrepancy measure. In a traditional VAE setting, the model is train to minimize a the sum of a reconstruction cost with the KL-divergence between the prior and the encoder acting as a regularizer.

In the case of WAE-GAN, the regularizer takes the form of a GAN-like discriminator function i.e. the regularizer can be express as $R_{WAE\text{GAN}}(\lambda, \gamma) = \lambda \log D_\gamma(\tilde{z})$ where D_λ is the discriminator parametrized by γ . The parameters of the discriminator are updated to maximized the default GAN objective function

$$(3) \quad E_{z \sim P_Z} \log D_\gamma(z) + E_{x \sim Q_{z|x}} \log(1 - D_\gamma(\tilde{z}))$$

The regularizer is maximized when both $D_\gamma(z)$ and $1 - D_\gamma(\tilde{z})$ have expected value of 0.5, i.e. when the discriminator can't distinguish sample's from the decoder from the one from the prior.

In the case of WAE-MMD, given a positive definite reproducing kernel $k(\cdot, \cdot)$ the penalty applied to Q_Z for being different from P_Z is defined as the norm of the difference between the expected value of $k(z, \cdot)$ over P_Z and the difference between the expected value of $k(z, \cdot)$ over Q_Z . Formal,

$$(4) \quad MMD_k(P_Z, Q_Z) = \left\| \int_{\mathcal{Z}} k(z, \cdot) p_Z(z) dz - \int_{\mathcal{Z}} k(z, \cdot) q_Z(z) dz \right\|_{\mathcal{H}_k}$$

where \mathcal{H}_k is the RKHS of real-valued functions mapping \mathcal{Z} to \mathbb{R} . Minimizing (4) ensure is equivalent to minimizing the point wise distance between the two expectation and thus it ensure that both distribution are close to each other. For computation purposes, an unbiased estimator of the MMD is used, i.e.

$$(5) \quad R_{WAE-MMD}(\lambda, k) = \frac{\lambda}{n(n-1)} \sum_{l \neq j} k(z_l, z_j) + \frac{\lambda}{n(n-1)} \sum_{l \neq j} k(\tilde{z}_l, \tilde{z}_j) - \frac{2\lambda}{n^2} \sum_{l,j} k(z_l, \tilde{z}_j)$$

3.3. Learning algorithm. Let D_X be the training set. Let ϕ and θ be the parameters of the decoder and the encoder respectively. The goal is to learn ϕ and θ that minimize the objective function (i.e. cost + regularizer). The following steps are repeated until convergence of ϕ and θ : (1) sample $\{x_1, \dots, x_n\} \in D_X$; (2) sample $z_1, \dots, z_n \sim P_z$; (3) sample $\tilde{z}_i \sim Q(Z|x_i)$, $i = 1, \dots, n$; (4) compute the gradient of the objective and update the parameters in the descending direction.

For the *WAE-GAN*, we also need to compute the gradient of the discriminator D_γ and update γ in the ascending direction (since we want to maximize (3)) before updating the objective.

4. EXPERIMENTS

4.1. Motivation for the experiments. To evaluate any improvement in learning the latent data distribution using the optimal transport-based Wasserstein distance over the standard variational auto encoder approach (making use of the KL-divergence), the authors trained 3 models: a 'stock' VAE, a WAE using the GAN-based Wasserstein penalty, and a WAE using the MDD-based Wasserstein penalty. The motivation was to empirically test whether the Wasserstein distance improves the reconstruction quality observed. The reconstruction analysis and latent space traversal analyses were standard. It was also good that they included a measure of image sharpness and the inception distance scores to quantify the claims made by this paper.

4.2. Reproducing the main results. The reproduction of this paper can be found in the following repo:

<https://github.com/josephdviviano/wae>

These experiments were run on both the celeba and mnist datasets. Note that, due to time restrictions, we only trained our models on 10% of the celeba dataset, and used a small test set of 1000 images.

Both the encoder-decoder architecture as well as the discriminator architecture (when applicable) used the ADAM optimizer. As in the original paper, the learning rate was halved at epoch

30, and then reduced further by a factor of 5 at epoch 50. We trained all models for 80 epochs. We used a convolutional architecture mirroring that used in the DC-GAN paper. The discriminator was a 4-layer fully connected network. Below, we report the hyperparameters used in each of our 6 analysis. All hyperparameters that are different from those in the original paper (out of necessity to stabilize training) are marked with a *:

NB: lambda refers to the scaling factor on the WAE penalty, scalebullshit refers to the scaling factor on the reconstruction loss and sigma refers to the variance of the latent layer used for sampled noise.

For the two WAE CELEBA runs, we used mean squared error reconstruction loss. For the VAE, we used binary cross entropy loss. The latent dimension was 64 for all analysis:

- **CELEBA VAE:** Batch size = 64, learning rate = 0.0001, lambda = N/A, scalebullshit = 0.05, sigma = 2.
- **CELEBA WAE-GAN:** Batch size = 64, learning rate = 0.0003, lambda = 1, scalebullshit = 0.05, sigma = 2.
- **CELEBA WAE-MMD:** Batch size = 64, learning rate = 0.00001*, lambda = 10*, scalebullshit = 1, sigma = 2.

Notes: For **CELEBA WAE-MMD**, we had to turn down lambda and the learning rate, and turn up scalebullshit, in order to stabilize training (i.e., more heavily weight the reconstruction part of the loss). Training of this model was quite unstable and required many attempts.

Surprisingly, we generally had more trouble stabilizing the training for the MNIST data. Again, the two WAE CELEBA runs used mean squared error reconstruction loss, while VAE run used binary cross entropy. The latent dimension was 8 for all analysis:

- **MNIST VAE:** Batch size = 100, learning rate = 0.001, lambda = N/A, scalebullshit = 1, sigma = 1*.
- **MNIST WAE-GAN:** Batch size = 100, learning rate = 0.001, lambda = 10, scalebullshit = 1, sigma = 2.
- **MNIST WAE-MDD:** Batch size = 100, learning rate = 0.00001*, scalebullshit = 1, sigma = 2.

Notes: **WAE-GAN** was relatively easy to train (with little instability issues). On the other hand **WAE-MMD** was extremely finicky, even worse than the CELEBA dataset. Gradients easily exploded, requiring very small learning rates.

Below we show the reconstructed data for all analysis at epoch 80. The authors also report the sharpness and inception scores for the CELEBA dataset, which we reproduce from our own results here:



FIGURE 2. Test CelebA reconstructions at epoch 80: original data (top left), VAE reconstructions (top right), WAE GAN reconstructions (bottom left), and WAE MDD reconstructions (bottom right).




FIGURE 3. Test MNIST reconstructions at epoch 80: using the VAE model.

TABLE 1. Blur and Fréchet Inception Distance (FID) Measures from CelebA dataset.

	Blur	FID
VAE	0.01947	0.1212
WAE-GAN	0.01151	0.1255
WAE-MMD	0.007089	0.1377

We note the following: **WAE-MMD** produces the lowest blur scores. In the original paper, **WAE-GAN** produced the lowest blur scores. However, both data types clearly produce less blurry images than **VAE**. In this sense we have replicated the performance advantage of the **WAE** family proposed by the authors. The results are less convincing for the Fréchet Inception Distance (FID) scores. The VAE actually outperforms both **WAE** measures, but the advantage isn't great and does not line up with the blur measure, nor a qualitative analysis of the images. Furthermore, the distance values are much smaller across the board than what was reported in the original paper, so we suspect there may have been an error in our analysis. Briefly, we used the first max-pooling features layer of the inception V3 network in all calculations. The mean and standard deviations from the activations from this layer were then used to compute the distance as:

$$DD^T + \text{trace}(\sigma_1) + \text{trace}(\sigma_2) - 2 * \sqrt{\sigma_1 \sigma_2^T}$$

Where $D = \mu_1 - \mu_2$, μ_x represents the mean activation from image x , and σ_x represents the standard deviation of activations from image x .

We suspect that we would have gotten better results if we had taken the activation from a lower layer of the inception V3 network. However, time constraints prevented us from re-running these analyses.

The blur measure is simply calculated as the variance of the output of the convolution of the reconstructed image with a laplacian filter.

Due to time constraints, we also did not replicate the latent-space traversal analysis. We do not consider it a main result because it does little to compare WAE with VAE, and we ran short on time due to having to debug the unstable training of the MNIST data and the MMD-based analysis. We acknowledge that this means we do not know for sure that the latent space learned by the WAE model is continuous (as would be expected in a variational model), but given that the other properties of the model were generally replicated, we assume this property would hold as well.

5. DISCUSSION

We are not critical of the experiments conducted considering the nature of the result. Namely, they present a new kind of generative model, and then compare the performance of this model using standard generative-model measures on stock datasets.

We are critical that not many quantitative results are shown. Training curves would have been welcome. There was also very little discussion of training stability, which we found to be an extremely challenging aspect of this replication.

We found multiple details of the implementation in the official source code not reported in the paper. All references below refer to the official code:

<https://github.com/tolstikhin/wae>

- For the reconstruction term of the loss (mean squared error), the authors **scale it by 5%** (line 335 of wae.py). This is obviously crucial, as it weights the reconstruction term of the WAE loss (`RECON + LAMBDA * PENALTY`) quite low compared to how they present the results in the paper. In our experiments, we tried using this scaling factor, lovingly called **SCALEBULLSHIT**.
- For the MMD loss, the MMD value is accumulated times over different scales (line 294 of wae.py). We copied this implementation detail, but are unsure of it's relevance. Nonetheless, this implementation detail should be covered in the paper, or at the very least, be justified by a comment in the code.
- We found training to be rather unstable for the MMD method in particular. Training tended to collapse if the mean or standard deviation of the encoded data grew too large, which happened frequently for the MMD method. The authors combat this by clamping the sigma value during reparameterization between -50 and 50, which they do not comment on anywhere in the paper or code. We found this to be a crucial implementation detail, although it did not combat the issue entirely.
- In general, the balance between the reconstruction term and the Wasserstein penalty is extremely sensitive. The authors would have done well to have given a theoretical intuition as to which this might be true as well as practical recommendations to mitigate this issue when trying to implement the WAE yourself.

Our major concern is the large difference in performance observed between using **WAE-GAN**, which we found to be both stable and produce very good results, and **WAE-MMD**, which we found to be extremely finicky and only sometimes produce reasonable results after fastidious hyper-parameter fine-tuning. The relationship between the reconstruction term and the WAE penalty is really evident when using the MMD method, but the authors don't give any explanation as to why either method would perform differently. Based on our experiments, we find it very hard to believe that the authors did not encounter training instabilities when using the MMD method, and would have liked to see the authors devote some attention to explaining the real-world differences between these two approaches to estimating the Wasserstein distance.

REFERENCES

[Tolstikhin et al.(2017)] Tolstikhin, I., Bousquet, O., Gelly, S., & Schoelkopf, B. 2017, arXiv:1711.01558

DÉPARTEMENT D'INFORMATIQUE ET DE RECHERCHE OPÉRATIONNELLE UNIVERSITÉ DE MONTRÉAL PAVILLON ANDRÉ-AISENSTADT, DIRO.